

Node.js Core 的启动过程及 V8 Snapshot 集成

张秋怡

Slides: <https://github.com/joyeecheung/talks>

Node 地下铁, 2019.09.08

概览

- ▶ 从 2018 年底开始进行重构启动过程代码
- ▶ **尽可能懒加载**
 - ▶ 不初始化不一定用到的东西
 - ▶ 重构后初始化依然需要加载至少 60 + 个内部 module

概览

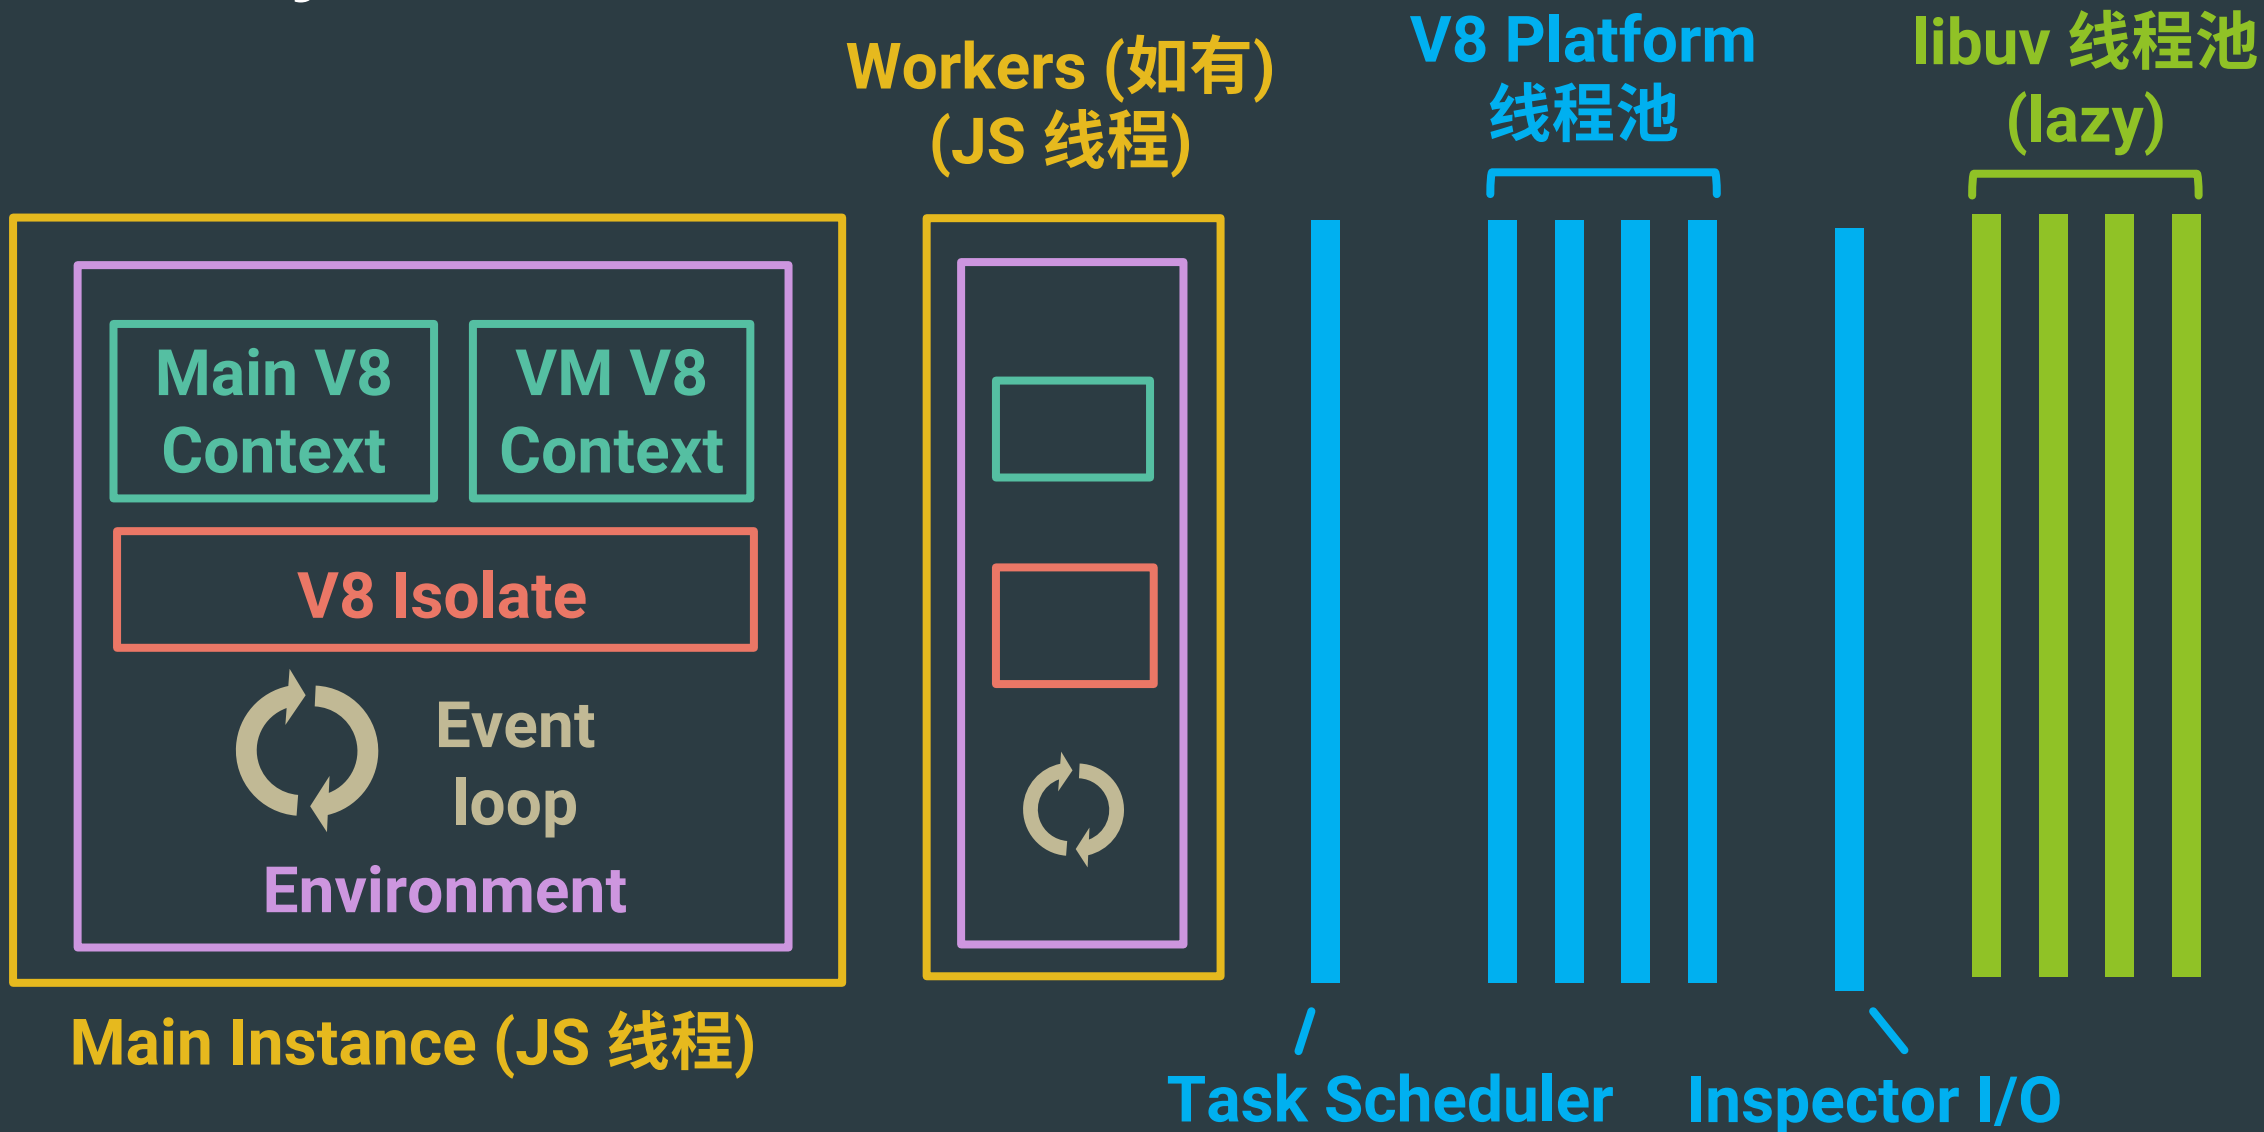
▶ 嵌入 code cache

- ▶ 包括编译出的 bytecode 等，预先编译好存储在 binary 内
- ▶ 启动时省去 parse 源代码生成 bytecode，但依然需要执行 bytecode 来完成初始化

▶ 嵌入 V8 Snapshot

- ▶ 不是 heap snapshot，是 startup snapshot
- ▶ 直接从预先编译生成的 snapshot 反序列化整个堆，不用执行代码来初始化到需要的状态

Node.js 进程模型



Bootstrap (2019.09)

```
node::Start()
```

Bootstrap (2019.09)

`node::Start()`



`InitializeOncePerProcess()`

Parse the CLI arguments, Initialize the V8 Platform, OpenSSL, ICU, signal handler...

Bootstrap (2019.09)

`node::Start()`



`InitializeOncePerProcess()`

Parse the CLI arguments, Initialize the V8 Platform, OpenSSL, ICU, signal handler...



`NodeMainInstance() / Worker()`

Bootstrap (2019.09)

node::Start()



InitializeOncePerProcess()

Parse the CLI arguments, Initialize the V8 Platform, OpenSSL, ICU, signal handler...



NodeMainInstance() / Worker() →

v8::Isolate

**JS heap, JS exceptions,
Microtask queue...**

Bootstrap (2019.09)

node::Start()



InitializeOncePerProcess()

Parse the CLI arguments, Initialize the V8 Platform, OpenSSL, ICU, signal handler...



NodeMainInstance() / Worker() →

v8::Isolate



v8::Context

**JS heap, JS exceptions,
Microtask queue...**

global proxy, JS builtins

Bootstrap (2019.09)

`node::Start()`



`InitializeOncePerProcess()`

Parse the CLI arguments, Initialize the V8 Platform, OpenSSL, ICU, signal handler...



`NodeMainInstance() / Worker()` →

`v8::Isolate`



`v8::Context`



`per_context/*.js`

**JS heap, JS exceptions,
Microtask queue...**

global proxy, JS builtins

Node.js primordials

Primordials

- ▶ JavaScript builtins 如 `Object`, `Object.prototype` 在启动时会被 clone 进一个对象, 并被 `Object.freeze()`, 用于给内部代码使用
- ▶ 为什么? 例: 用户可以 `delete Function.prototype.call`
- ▶ 逐步迁移所有内部代码到使用这些 primordials

```
~/projects/node ➤ master ➤ out/Release/node
```

```
Welcome to Node.js v13.0.0-pre.
```

```
Type ".help" for more information.
```

```
> delete Function.prototype.call
```

```
Thrown:
```

```
TypeError: _memory.call is not a function
```

```
    at finish (repl.js:721:15)
```

```
    at finishExecution (repl.js:379:7)
```

```
    at REPLServer.defaultEval (repl.js:465:7)
```

```
    at bound (domain.js:423:14)
```

```
    at REPLServer.runBound [as eval] (domain.js:436:12)
```

Bootstrap (2019.09)

node::Start()



InitializeOncePerProcess()

Parse the CLI arguments, Initialize the V8 Platform, signal handler...



NodeMainInstance() / Worker() →

v8::Isolate



v8::Context



per_context/*.js



node::Environment

JS heap, JS exceptions,
Microtask queue...

global proxy, JS builtins

Node.js primordials

不知道放哪里好就放在
这里的各种数据

Bootstrap (2019.09)

node::Start()



InitializeOncePerProcess()

Parse the CLI arguments, Initialize the V8 Platform, signal handler...



NodeMainInstance() / Worker() →

v8::Isolate



v8::Context



per_context/*.js



node::Environment



libuv handles

JS heap, JS exceptions,
Microtask queue...

global proxy, JS builtins

Node.js primordials

不知道放哪里好就放在
这里的各种数据

Bootstrap (2019.09)

node::Start()



InitializeOncePerProcess()

Parse the CLI arguments, Initialize the V8 Platform, signal handler...



NodeMainInstance() / Worker() →

v8::Isolate



v8::Context



per_context/*.js



node::Environment



libuv handles



inspector agent

JS heap, JS exceptions,
Microtask queue...

global proxy, JS builtins

Node.js primordials

不知道放哪里好就放在
这里的各种数据

Bootstrap (2019.09)

`node::Start()`



`InitializeOncePerProcess()`

Parse the CLI arguments, Initialize the V8 Platform, signal handler...



`NodeMainInstance()` / `Worker()` →

`v8::Isolate`



`v8::Context`



`per_context/*.js`



`node::Environment`



libuv handles



inspector agent



`bootstrap/*.js`

JS heap, JS exceptions,
Microtask queue...

global proxy, JS builtins

Node.js primordials

不知道放哪里好就放在
这里的各种数据

global, process, task
queues, ESM/CJS
loaders ...

lib/internal/bootstrap/loaders.js

- ▶ Internal module loaders
- ▶ **C++** binding loaders
 - ▶ `process.binding()`
 - ▶ `process._linkedBinding()`
 - ▶ `internalBinding()`
- ▶ `require()` for loading other internal **JavaScript** modules

Built-in Modules (Native Modules)

lib/*.js

```
"use strict";  
...
```

JavaScript code

tools/js2c.py

NativeModuleLoader::LoadJavaScriptSource()

```
static const uint16_t timers_raw[] = {  
    39, 117, 115, 101...  
};
```

static data array 包含
builtins 的源代码

Built-in Modules (Native Modules)

lib/*.js

```
"use strict";  
...
```

JavaScript code

tools/mkcodecache

NativeModuleEnv::InitializeCodeCache()

```
static const uint8_t assert[] = {  
    165, 3, 222, 192, 132, ...  
};
```

static data array 包含
code cache

Built-in Modules (Native Modules)

```
function (exports, require, module, process,  
        internalBinding, primordials) {  
  
    require('internal/fs/utils');  
  
    module.exports = {...};  
  
}
```

Compiled with a special wrapper
that include access to more internals

lib/internal/bootstrap/node.js

- ▶ 初始化 `process` 和 `global` 上面的大部分成员
- ▶ C++ 将 `isMainThread`, `ownsProcessState` 传入，用于针对不同的场景进行不同的初始化
 - ▶ worker 里是 `false`, 主线程里是 `true`

lib/internal/bootstrap/node.js

- ▶ 初始化将会作为 `v8::Persistent` 存储在 `Environment` 的 JavaScript callbacks
 - ▶ Async hook callbacks
 - ▶ Timers & `process.nextTick()` schedulers
- ▶ 不能进行任何异步操作 (无法从 snapshot 还原)
- ▶ 不能依赖任何命令行参数和环境变量。如果打包进 snapshot , 编译环境 (Node.js 发布用的集群) 和运行环境 (用户机器) 不一致会导致大部分参数无法生效。

lib/internal/bootstrap/pre_execution.js

- ▶ 被 main scripts 加载（见后），不主动运行
- ▶ 主要根据命令行参数和环境变量进行进一步的初始化
 - ▶ 包括 CJS & ESM loaders
- ▶ 不包含在 snapshot

```
if (!getOptionValue('--no-warnings') &&  
    process.env.NODE_NO_WARNINGS !== '1') {  
  process.on('warning', onWarning);  
}
```

User land CommonJS Modules

► Loader 实现在 `lib/internal/modules/cjs/`

```
function (exports, require, module, __filename, __dirname) {  
    require('fs');  
}
```

Wrap user code with objects initialized by Node.js

Built-in Modules (Native Modules)

```
function (exports, require, module, process,  
        internalBinding, primordials) {  
  
    require('internal/fs/utils');  
  
    module.exports = {...};  
  
}
```

Compiled with a special wrapper
that include access to more internals

User land ECMAScript Modules

- ▶ Loader 实现在 `lib/internal/modules/esm/`
- ▶ 不改变 context, 只有 bootstrap scripts 往 global proxy 注入的全局变量 (这部分不分 ESM 和 CJS)
 - ▶ `Buffer`, `process`, etc.

User land ECMAScript Modules

- ▶ 一个内部的 `WeakMap` 包含 `ModuleWrap` -> `Options`
 - ▶ `Options` 包括 dynamic `import()` 的 callback 和 `import.meta` data
 - ▶ Per-isolate
 - ▶ `HostImportModuleDynamicallyCallback`
 - ▶ `HostInitializeImportMetaObjectCallback`

Bootstrap (2019.09)

`node::Start()`



`InitializeOncePerProcess()`

Parse the CLI arguments, Initialize the V8 Platform, signal handler...



`NodeMainInstance()` / `Worker()` →

`v8::Isolate`
↓
`v8::Context`
↓
`per_context/*.js`
↓
`node::Environment`
↓
libuv handles
↓
inspector agent
↓
`bootstrap/*.js`
↓
`main/? .js`

JS heap, JS exceptions,
Microtask queue...

global proxy, JS builtins

Node.js primordials

不知道放哪里好就放在
这里的各种数据

global, process, task
queues, ESM/CJS
loaders ...

e.g.
`run_main_module.js`

Main scripts

- ▶ `lib/internal/main/*.js`
- ▶ 主线程
 - ▶ `StartMainThreadExecution()`
 - ▶ 根据命令行参数等条件，选择一个 main script 运行
- ▶ Workers
 - ▶ `worker_thread.js`
- ▶ 先加载 `lib/internal/bootstrap/pre_execution.js`
根据运行环境进行初始化

Main scripts

- ▶ `check_syntax.js`: `node -c test.js`
- ▶ `eval_stdin.js`: `cat test.js | node -e`
- ▶ `eval_string.js`: `node -e '1'`
- ▶ `inspect.js`: `node inspect ...`
- ▶ `print_bash_completion.js`: `node --completion-bash`
- ▶ `print_help.js`: `node --help`
- ▶ `prof_process.js`: `node --prof-process v8.log`

Main scripts

- ▶ `run_third_party_main.js`
 - ▶ 运行 embedders 嵌入的 `lib/_third_party_main.js`
- ▶ `environment.js`
 - ▶ For C++ test fixtures

Requested features

- ▶ Customize entry point for bundled CLI tools
- ▶ Better entry points for embedders

Main scripts

- ▶ repl.js: node
- ▶ worker_thread.js: for workers
- ▶ run_main_module.js
 - ▶ node index.js
 - ▶ node --experimental-modules index.mjs

Main scripts

- ▶ repl.js: node
- ▶ run_main_module.js
 - ▶ node index.js
 - ▶ node --experimental-modules index.mjs
- ▶ worker_thread.js: for workers

Bootstrap (2019.09)

`node::Start()`



`InitializeOncePerProcess()`

Parse the CLI arguments, Initialize the V8 Platform, signal handler...

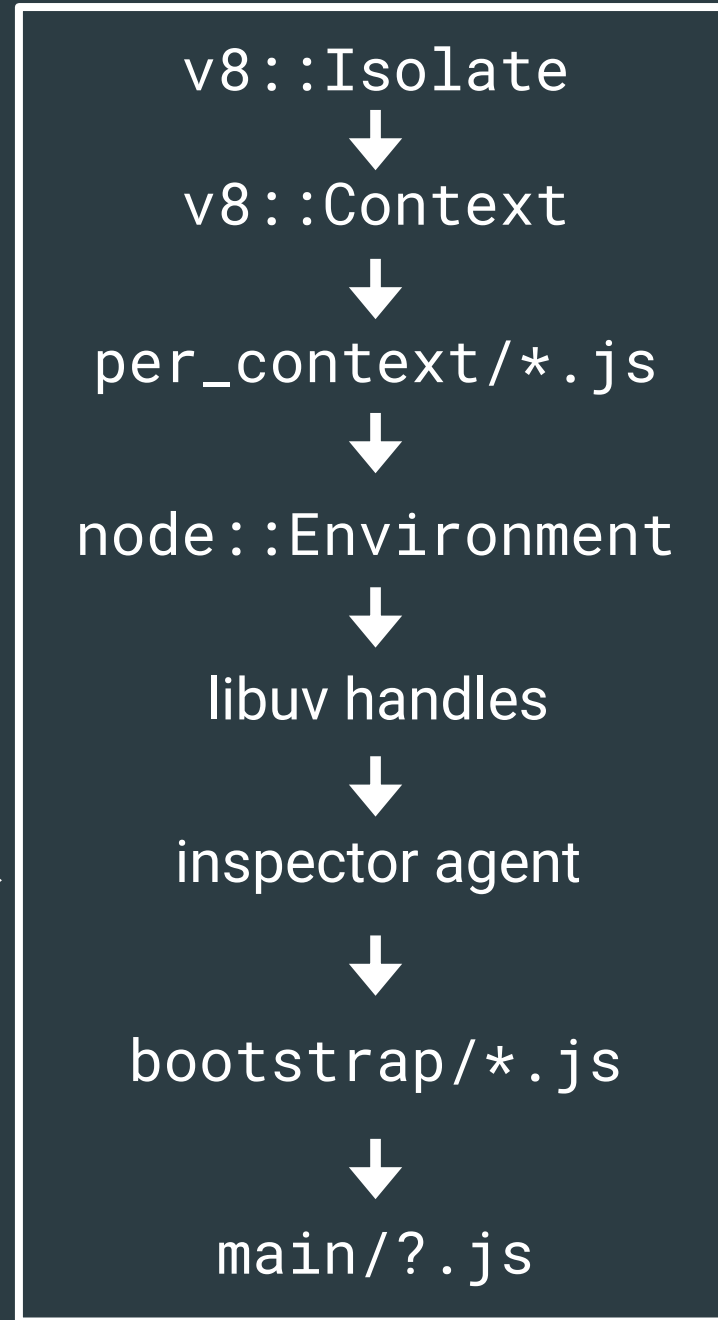


`NodeMainInstance() / Worker()` →



```
do {  
  uv_run(...)  
} while (...)
```

Event Loop



JS heap, JS exceptions, Microtask queue...

global proxy, JS builtins

Node.js primordials

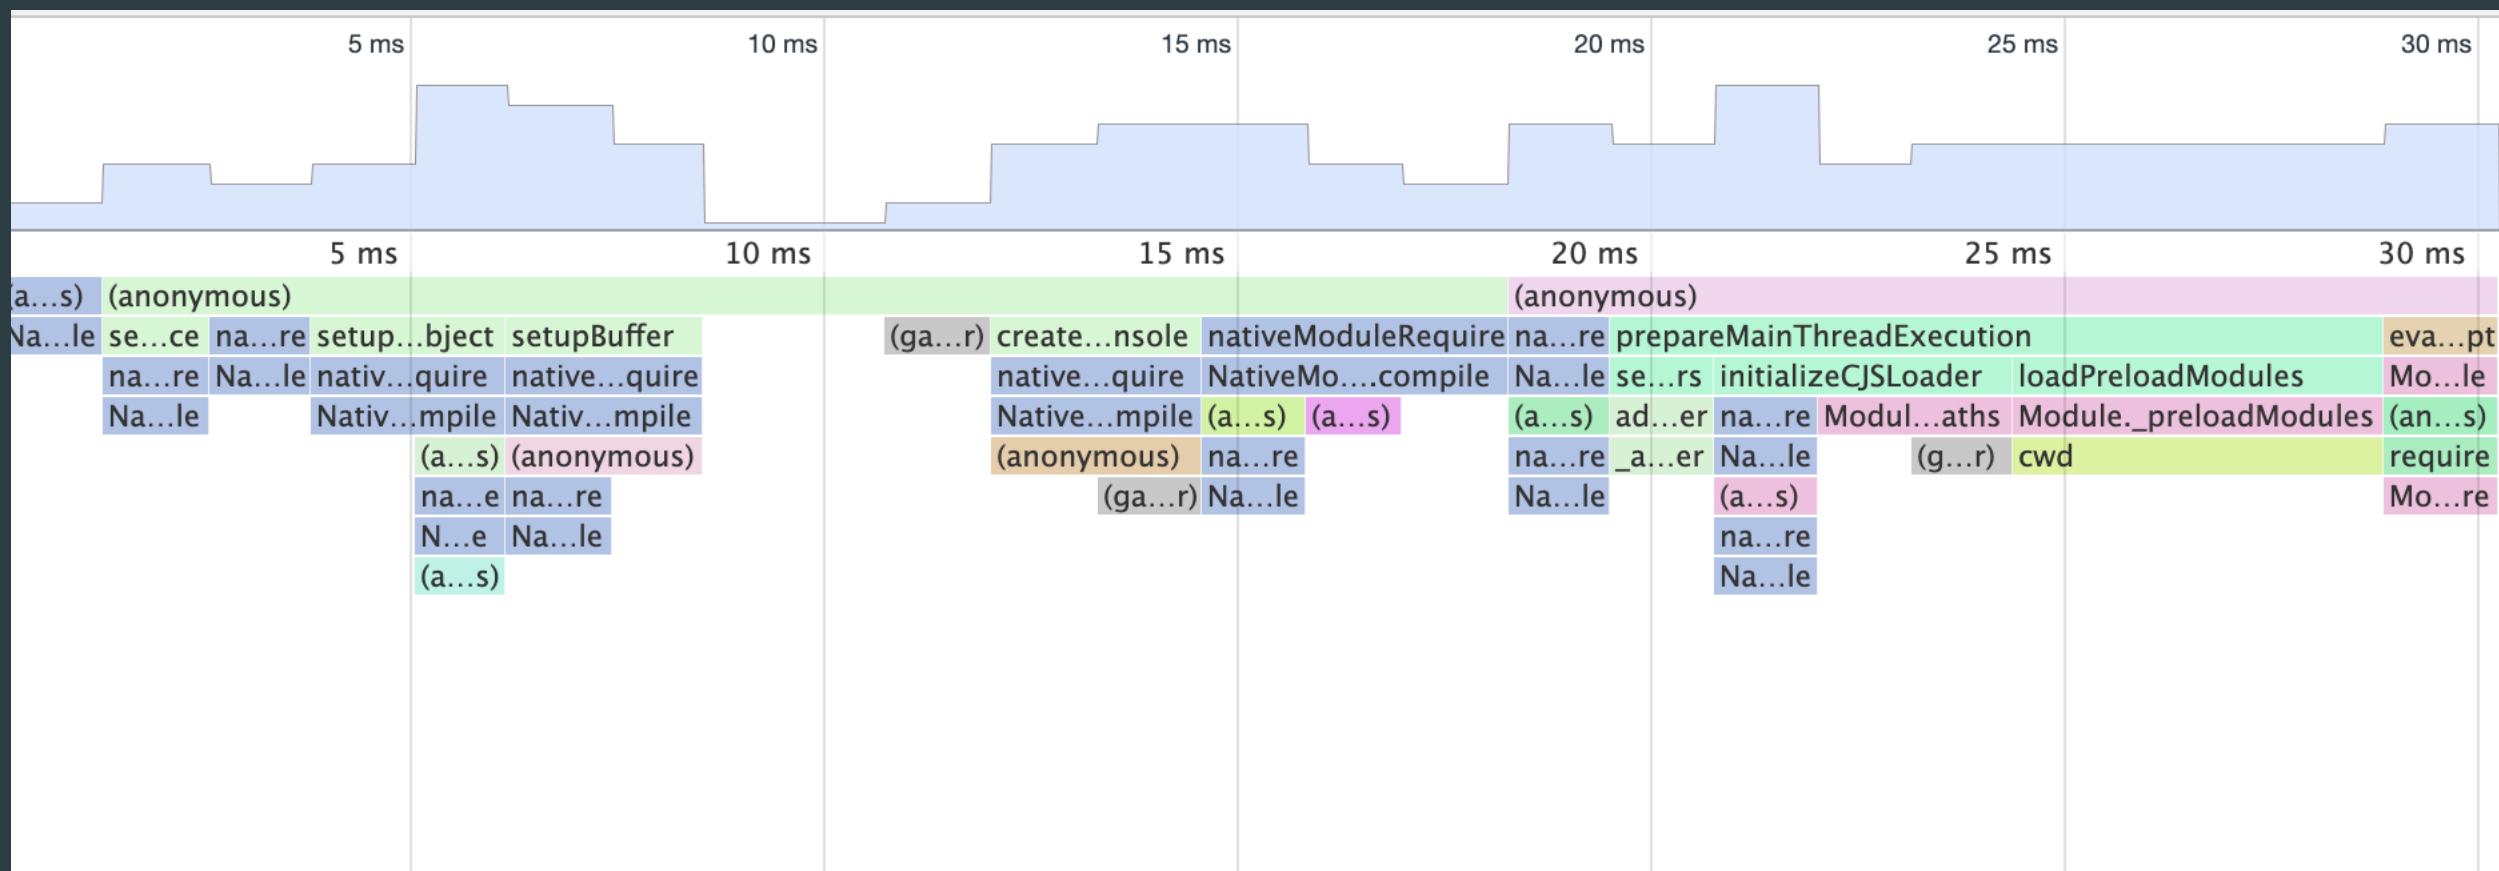
不知道放哪里好就放在这里的各种数据

global, process, task queues, ESM/CJS loaders ...

e.g. run_main_module.js

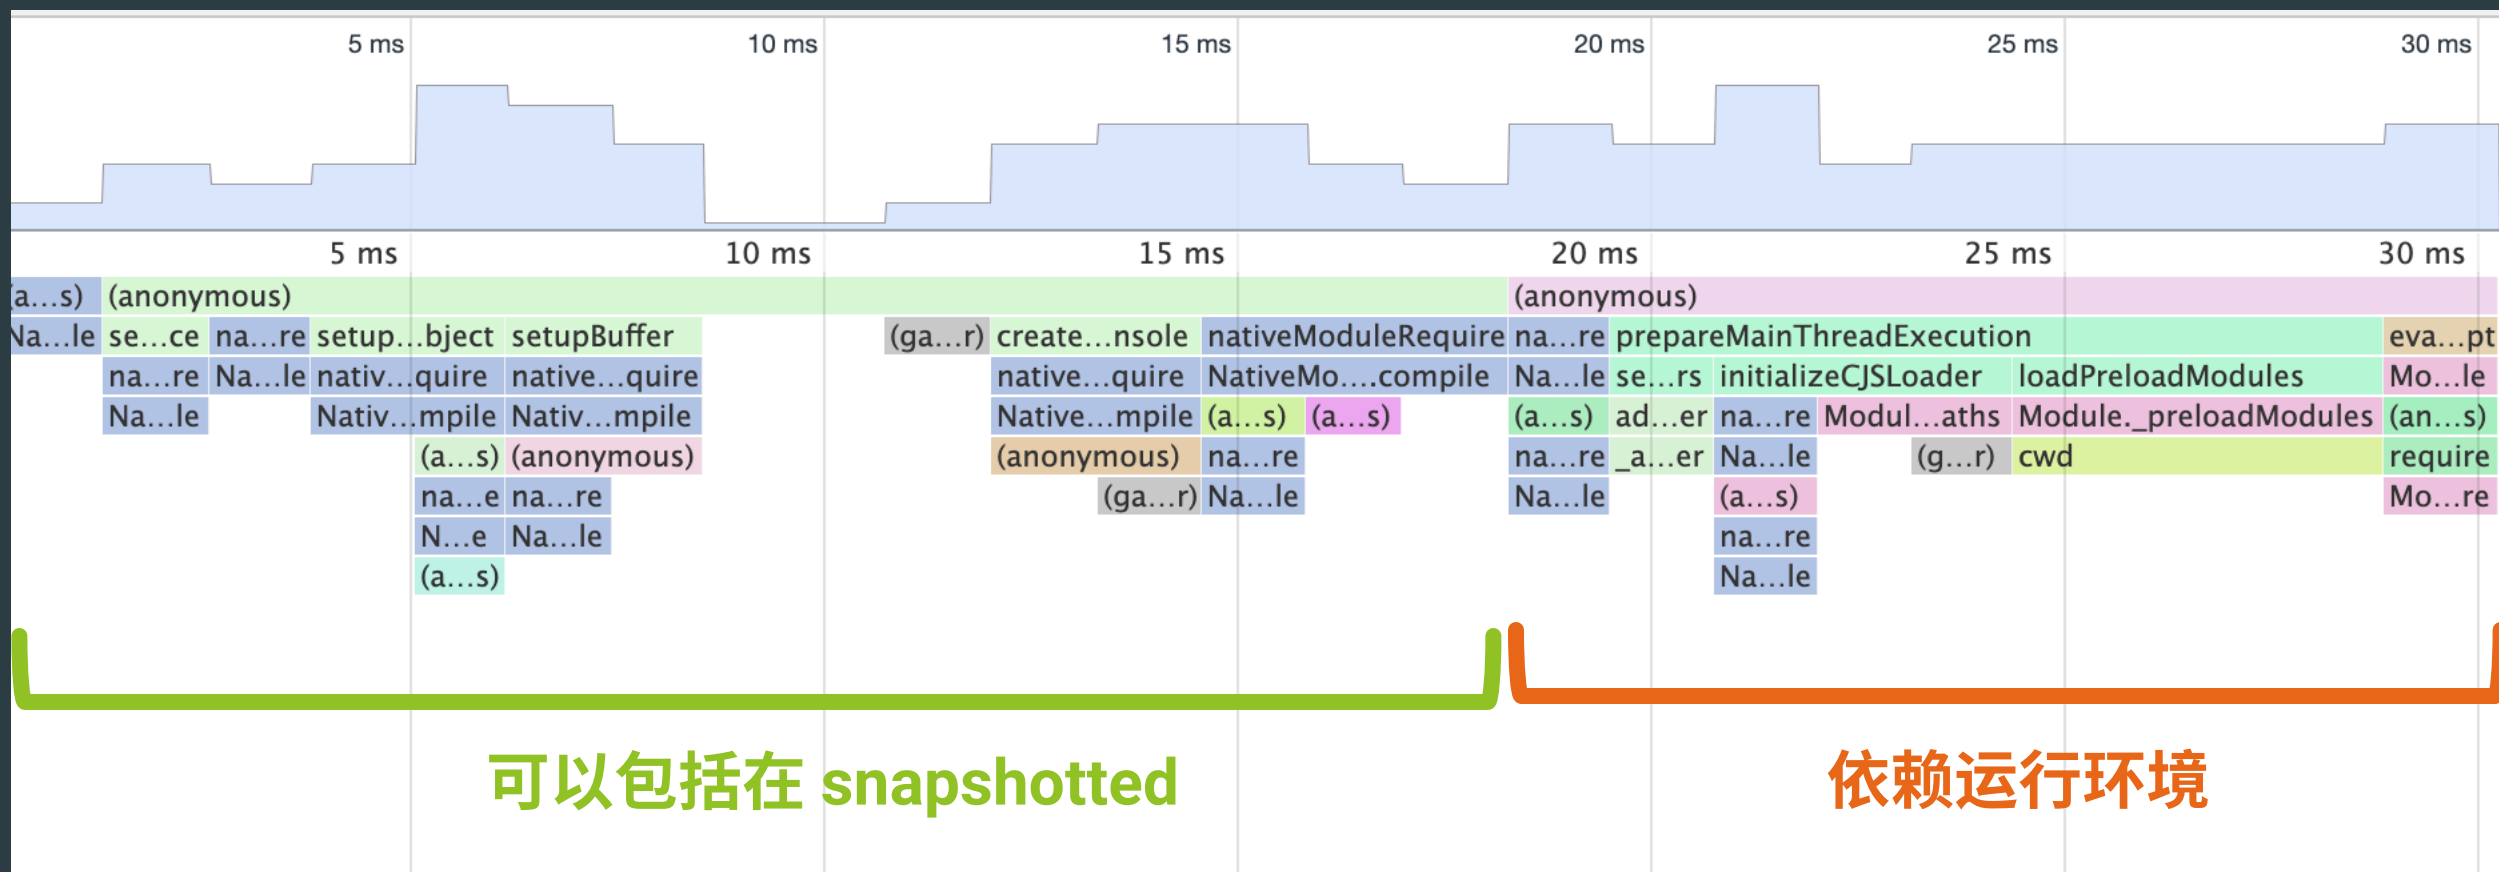
Snapshot 集成

```
out/Release/node --cpu-prof-interval=100 --cpu-prof -e "{}"
```



Snapshot 集成

```
out/Release/node --cpu-prof-interval=100 --cpu-prof -e "{}"
```



Snapshot 集成

```
bash-5.0$ time luajit -e "local x = 1"
```

```
real    0m0.005s
```

```
user    0m0.002s
```

```
sys     0m0.002s
```

```
bash-5.0$ time perl -e 1
```

```
real    0m0.007s
```

```
user    0m0.003s
```

```
sys     0m0.003s
```

```
bash-5.0$ time ~/.jsvu/v8 -e 1
```

```
real    0m0.024s
```

```
user    0m0.009s
```

```
sys     0m0.012s
```

```
bash-5.0$ time out/Release/node -e 1
```

```
real    0m0.038s
```

```
user    0m0.028s
```

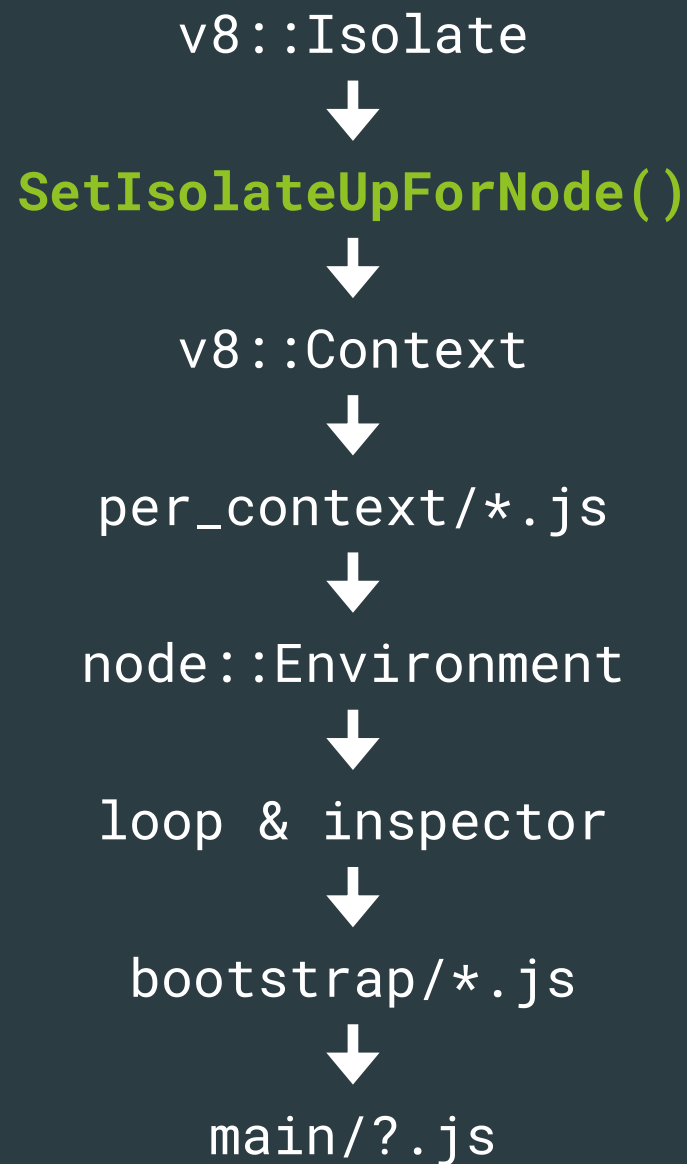
```
sys     0m0.007s
```

d8 with default snapshot

node master without snapshot

Snapshot 集成

Original



Snapshot 集成

Snapshotted (2019.05)

v8::Isolate



Context::FromSnapshot()



SetIsolateUpForNode()

Re-install callbacks



node::Environment



loop & inspector



bootstrap/*.js



main/? .js

Snapshot 集成

目标

直接从 snapshot 加载部分初始化好的环境，
而不是从头执行 `per_context/*.js`
& `bootstrap/*.js` 来初始化

`v8::Isolate`
↓
`Context::FromSnapshot()`
↓
`Environment::FromSnapshot()`
↓

`SetIsolateUpForNode()`
Re-install callbacks

↓
`loop & inspector`

↓
`main/? .js`

Snapshot 集成

重构

在截取 snapshot 前的启动流程必须不能依赖运行环境

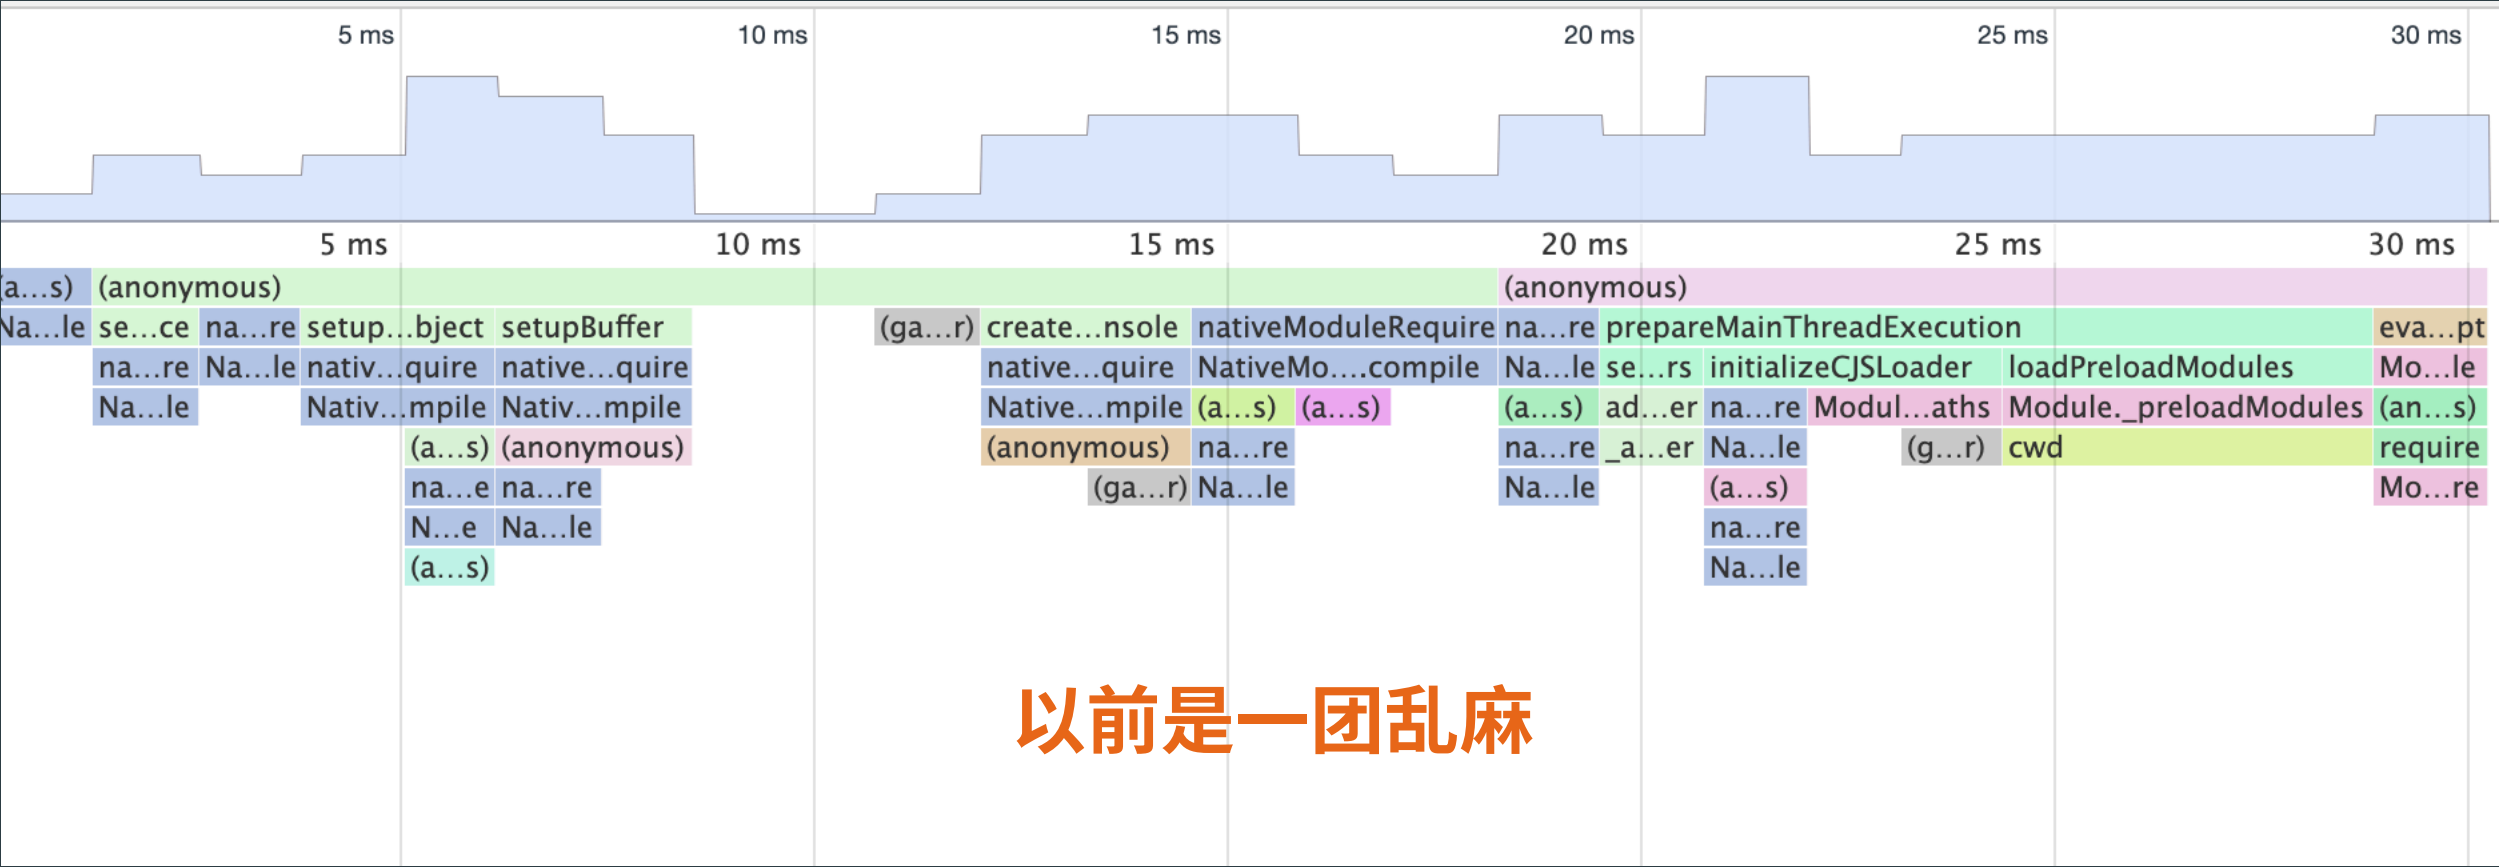
v8::Isolate
↓
Context::FromSnapshot()
↓
Environment:: FromSnapshot()
↓

SetIsolateUpForNode()
Re-install callbacks

↓
loop & inspector

↓
main/? .js

Snapshot 集成



以前是一团乱麻

Snapshot 集成

```
if (!getOptionValue('--no-warnings') &&  
    process.env.NODE_NO_WARNINGS !== '1') {  
  process.on('warning', onWarning);  
}
```

重构

lib/internal/bootstrap/pre_execution.js

v8::Isolate



Context::FromSnapshot()



Environment::FromSnapshot()



SetIsolateUpForNode()

Re-install callbacks



loop & inspector



main/? .js

Snapshot 集成

重构

调整步骤，需要在 snapshot 加载后重新同步 C++ 部分的状态 (snapshot 里只保存了 JavaScript 的状态)

```
v8::Isolate  
↓  
Context::FromSnapshot()  
↓  
Environment:: FromSnapshot()  
↓
```

```
SetIsolateUpForNode()  
Re-install callbacks
```

```
↓  
loop & inspector
```

```
↓  
main/? .js
```

目前进展

- ▶ Per-context scripts are snapshotted and shipped by default since v12.5.0
- ▶ **v12.5.0 v.s. v11.2.0**
 - ▶ ~60% faster child process startup
 - ▶ ~120% faster worker startup
- ▶ 大部分提升来自重构（更多懒加载）和 embedded code cache

进行中的工作

- ▶ V8 issues
 - ▶ Rehashing Map & Set
 - ▶ 优化 `v8::External`
 - ▶ Lazy initialization of ICU
- ▶ V8 部分解决后, `bootstrap/loaders.js` 和 `bootstrap/node.js` 的运行结果可以包含入 `embedded snapshot`
 - ▶ <https://github.com/nodejs/node/issues/17058> 中的初步实现提升为 4x
 - ▶ 实际可能低一些, 因为这个原型包含了一些不应该被截进去的状态

Future plans

- ▶ User-land snapshot builder & loader
 - ▶ 改进开发过程中的错误提示
 - ▶ 遍历到 context-dependent 的对象无法序列化
 - ▶ 遇到未知外部引用
 - ▶ 如何整合 C++ addons?
 - ▶ Warm up
 - ▶ 对包含在 snapshot 内的运行状态 (e.g. 环境变量) 进行提示
 - ▶ 用户自己打包的应用可以自行保证编译环境和运行环境一致

Thank you